

# Bayesian ridge regression and Bayesian lasso

Frank van der Meulen (TU Delft)

1 May 2015


## ILLUSTRATION OF THE FULL RIDGE AND LASSO REGULARISATION PATHS USING GLMNET

The glmnet package contains many function for estimation of generalised linear models using  $\ell_1$  and or  $\ell_2$  regularisation. We illustrate it with a dataset that is available within R.


```
library(glmnet)
library(lars) # contains the diabetes data
data(diabetes)
d <- diabetes

lasfit <- glmnet(d$x,d$y,alpha=1) # lasso
ridfit <- glmnet(d$x,d$y,alpha=0) # ridge


plot(ridfit,lwd=2,xvar='lambda',label=TRUE)
abline(0,0)
```



```
plot(lasfit,lwd=2,xvar='lambda',label=TRUE)
abline(0,0)
```



```
# cross validation plot for lasso
cvfit = cv.glmnet(d$x, d$y, alpha=1, nfolds=20) # 5-fold cv
plot(cvfit)
```



```
cvfit$lambda.min  # value of lambda that gives minimum cvm.
```

```
## [1] 1.092931
```

```
cvfit$lambda.1se  # largest value of lambda such that error is within 1 standard error of the minimum.
```

```
## [1] 5.314486
```

```
coef(lasfit,s=cvfit$lambda.min)
```

```
## 11 x 1 sparse Matrix of class "dgCMatrix"
##              1
## (Intercept) 152.13348
## age          .
## sex          -192.47489
## bmi          521.64423
## map          294.40618
## tc           -97.72765
## ldl          .
## hdl          -222.10137
## tch          .
## ltg          511.10045
## glu          52.15790
```

```
coef(lasfit,s=cvfit$lambda.1se)
```

```
## 11 x 1 sparse Matrix of class "dgCMatrix"
##              1
## (Intercept) 152.13348
## age          .
## sex          -33.35229
## bmi          508.13935
## map          210.34606
## tc           .
## ldl          .
## hdl          -138.84433
## tch          .
## ltg          444.59064
## glu          .
```

## ILLUSTRATION OF RIDGE, LASSO, BAYESIAN RIDGE AND BAYESIAN LASSO

Install the monomvn package and call the library.

```
library(monomvn)
```

The following function generates sparse data from the linear regression model:

```

genSparseData <- function(n,p,beta0.nonzero,sd.noise)
# generates the n x p design matrix with all elements drawn from a standard normal distribution
# beta0 is the vector beta0.nonzero, augmented with zeros
# a list containing (x, y, beta0) is returned, where
# y = x * beta0 + eps,
# where eps is a vector of independent normal random variables
# with mean zero and st.dev equal to sd.noise
{
  r <- length(beta0.nonzero)
  if (p< (r+2))
  { stop('p is too small in relation with the nr of nonzero elements in beta0.nonzero')
  } else {
    x <- matrix(rnorm(n*p),nrow=n)
    beta0 <- c(beta0.nonzero, rep(0,p-r))
    y <- x %*% beta0 + rnorm(n, sd=sd.noise)
    list(x=x,y=y, beta0=beta0)
  }
}

```

## TEST CASE: n=25, p=500

Generate the data

```

set.seed(38) # fix the seed of the RNG
d1 <- genSparseData(25,100,c(5, -3, 1, 0.2),.1)
str(d1)

```

```

## List of 3
## $ x      : num [1:25, 1:100] -0.2536 -1.0556 0.6865 0.0252 -1.6718 ...
## $ y      : num [1:25, 1] -1.54 -4.36 4.97 -3.39 -7.52 ...
## $ beta0: num [1:100] 5 -3 1 0.2 0 0 0 0 0 0 ...

```

## Ordinary Least Squares regression

As  $p > n$  this will not work:

```

reg.ols <- lm(d1$y ~ d1$x)
summary(reg.ols)$coef[1:50] # print first 50 coefs

```

```

## [1] -0.283462782  4.669685040 -2.918809552  0.934978200 -0.150529666
## [6]  0.593930135 -0.249878902 -0.364522097  0.020569314 -0.090580776
## [11]  0.031054241 -0.180278334  0.417571414  0.410551748  0.108340973
## [16]  0.186894642 -0.006294776 -0.124836427 -0.321128549  0.396752935
## [21] -0.599354235  0.179711848 -0.245206735 -0.110487848  0.251961652
## [26]          NaN          NaN          NaN          NaN          NaN
## [31]          NaN          NaN          NaN          NaN          NaN
## [36]          NaN          NaN          NaN          NaN          NaN
## [41]          NaN          NaN          NaN          NaN          NaN
## [46]          NaN          NaN          NaN          NaN          NaN

```

## Ridge regression

The default choice for setting the regularisation parameter is 10-fold cross-validation.

```
rid <- regress(d1$x, d1$y, method="ridge")
# note that this function itself always automatically includes an intercept

ridCoef <- rid$b      # vector of coefficients (note that an intercept is included!)
ridCoef[1:15]         # print first 15 elements
```

```
## [1] -0.15736762  0.92029473 -0.53315934  0.10373099  0.13218585
## [6]  0.32623143  0.21728364 -0.03910456  0.25160404 -0.02652536
## [11]  0.15573363 -0.15967645 -0.03379762  0.19125742 -0.03205427
```

```
rid$lambda           # regularisation parameter
```

```
## [1] 4.317531
```


## Bayesian ridge regression

```
ridBayes <- bridge(d1$x, d1$y, T=1000)
```

```
## t=100, m=23
## t=200, m=24
## t=300, m=24
## t=400, m=23
## t=500, m=6
## t=600, m=4
## t=700, m=7
## t=800, m=4
## t=900, m=7
```

```
BI <- 250 # choose a value for burnin
plot(ridBayes, burnin=BI)
```

## Boxplots of regression coefficients



```
ridBayesCoef <- cbind(ridBayes$mu, ridBayes$beta)[BI:T,]
ridgeBayesMean <- colMeans(ridBayesCoef)
ridgeBayesMedian <- apply(ridBayesCoef,2,median)
results <- data.frame(truth=c(0,d1$beta0), ridge=rid$b,
                      Bridge.mean=ridgeBayesMean,
                      Bridge.median=ridgeBayesMedian)
results[1:15,]
```

##	truth	ridge	Bridge.mean	Bridge.median
##	0.0	-0.15736762	-0.078330593	-0.01617783
## b.1	5.0	0.92029473	2.325002119	2.47141466
## b.2	-3.0	-0.53315934	-1.420499366	-1.54504220
## b.3	1.0	0.10373099	0.030292466	0.00000000
## b.4	0.2	0.13218585	0.024359859	0.00000000
## b.5	0.0	0.32623143	0.046947740	0.00000000
## b.6	0.0	0.21728364	0.012758534	0.00000000
## b.7	0.0	-0.03910456	-0.009934238	0.00000000
## b.8	0.0	0.25160404	0.027456727	0.00000000
## b.9	0.0	-0.02652536	-0.011492532	0.00000000
## b.10	0.0	0.15573363	0.129380055	0.00000000
## b.11	0.0	-0.15967645	-0.005949868	0.00000000
## b.12	0.0	-0.03379762	-0.002303823	0.00000000
## b.13	0.0	0.19125742	0.088334363	0.00000000
## b.14	0.0	-0.03205427	0.009160131	0.00000000

Here we have taken a fairly small number of iterations, if possible take more in real applications.

Many results can be obtained from the summary function

```
?summary.blasso
```

For example, the estimated posterior probability that the individual components of the regression coefficients  $\beta$  is nonzero is obtained from

```
bn0 <- summary(ridBayes)$bn0
bn0[1:15]    # print the first 15 elements
```

```
##   b.1   b.2   b.3   b.4   b.5   b.6   b.7   b.8   b.9  b.10  b.11  b.12
## 0.975 0.979 0.618 0.631 0.104 0.102 0.071 0.078 0.088 0.105 0.116 0.070
##   b.13  b.14  b.15
## 0.178 0.161 0.115
```

## Lasso regression

```
las <- regress(d1$x, d1$y, method="lasso")
lasCoef <- las$b
lasCoef[1:15]
```

```
## [1] -0.03637022  4.90330600 -2.96080008  0.96275693  0.12709471
## [6]  0.00000000  0.00000000  0.00000000  0.00000000  0.00000000
## [11] 0.00000000  0.00000000  0.00000000  0.00000000  0.00938084
```

```
lasLambda <- las$lambda
lasLambda
```

```
## [1] 0.1271622
```


## Bayesian lasso regression

```
lasBayes <- blasso(d1$x, d1$y, T=1000)
```

```
## t=100, m=14
## t=200, m=11
## t=300, m=10
## t=400, m=19
## t=500, m=10
## t=600, m=24
## t=700, m=21
## t=800, m=8
## t=900, m=10
```

```
BI <- 250    # choose a value for burnin
plot(lasBayes, burnin=BI)
```

## Boxplots of regression coefficients



```
lasBayesCoef <- cbind(lasBayes$mu, lasBayes$beta)[BI:T,]
lassoBayesMean <- colMeans(lasBayesCoef)
lassoBayesMedian <- apply(lasBayesCoef,2,median)
results <- data.frame(truth=c(0,d1$beta0), lasso=las$b,
                      Blasso.mean=lassoBayesMean,
                      Blasso.median=lassoBayesMedian)
results[1:15,]
```

##	truth	lasso	Blasso.mean	Blasso.median
##	0.0	-0.03637022	-5.528043e-02	-0.03030157
## b.1	5.0	4.90330600	4.714824e+00	4.96114620
## b.2	-3.0	-2.96080008	-2.841021e+00	-2.99477934
## b.3	1.0	0.96275693	9.482718e-01	1.00673329
## b.4	0.2	0.12709471	1.159496e-01	0.15135666
## b.5	0.0	0.00000000	-7.536528e-04	0.00000000
## b.6	0.0	0.00000000	-2.281772e-03	0.00000000
## b.7	0.0	0.00000000	5.834138e-06	0.00000000
## b.8	0.0	0.00000000	-4.407082e-04	0.00000000
## b.9	0.0	0.00000000	5.904413e-04	0.00000000
## b.10	0.0	0.00000000	-3.278002e-05	0.00000000
## b.11	0.0	0.00000000	6.732857e-03	0.00000000
## b.12	0.0	0.00000000	-7.400661e-04	0.00000000
## b.13	0.0	0.00000000	1.134870e-02	0.00000000
## b.14	0.0	0.00938084	1.564100e-02	0.00000000

## Horseshoe regression


```
hs <- bhs(d1$x,d1$y,T=1000)
```



```
## t=100, m=24
## t=200, m=24
## t=300, m=22
## t=400, m=24
## t=500, m=24
## t=600, m=24
## t=700, m=24
## t=800, m=24
## t=900, m=24
```

```
BI <- 250 # choose a value for burnin
plot(hs, burnin=BI)
```

## Boxplots of regression coefficients



```
hsCoef <- cbind(hs$mu, hs$beta)[BI:T,]
hsMean <- colMeans(hsCoef)
hsMedian <- apply(hsCoef, 2, median)
results <- data.frame(truth=c(0, d1$beta0),
                      hs.mean=hsMean,
                      hs.median=hsMedian)
results[1:15,]
```

##	truth	hs.mean	hs.median
##	0.0	0.0390685009	-0.03415112
## b.1	5.0	4.6854096620	4.95795955
## b.2	-3.0	-2.4018622091	-2.99233925
## b.3	1.0	0.6810286939	0.99902745
## b.4	0.2	0.1101165560	0.14864912
## b.5	0.0	0.0083831252	0.00000000
## b.6	0.0	0.0006003356	0.00000000
## b.7	0.0	-0.0112578890	0.00000000
## b.8	0.0	-0.0107512720	0.00000000



```
## b.9      0.0 -0.0009990126  0.00000000
## b.10     0.0  0.0200152191  0.00000000
## b.11     0.0 -0.0059784126  0.00000000
## b.12     0.0  0.0031806283  0.00000000
## b.13     0.0  0.0014306550  0.00000000
## b.14     0.0  0.0068611827  0.00000000
```

If you wish to make additional plots, consult the help page for the plotting function

```
?plot.blasso
```

As an example, suppose for the Bayesian lasso we wish to make a traceplot for lambda:

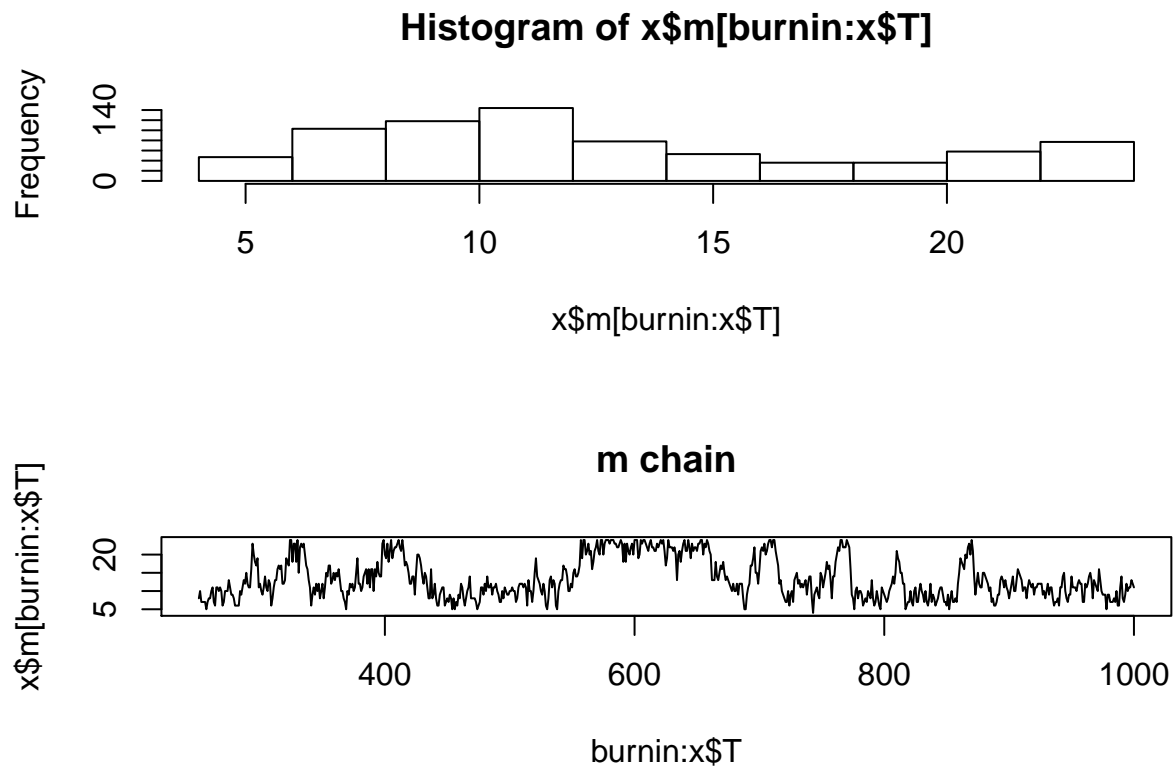
```
plot(lasBayes,burnin=BI,which='lambda')
```



Or

plot the size of different models visited

```
plot(lasBayes,burnin=BI,which='m')
```



## ILLUSTRATION OF RECOVERING A SPARSE SIGNAL

Here is another illustration where we try to recover a sparse signal. We use both the (frequentist) lasso and horseshoe

```
IT <- 1000
BI <- 250

# truth
n <- 10
y0 <- rep(0,n^2)
y0[13:15] <- c(3,-2,5)
y0[65] <- 3

# observations
y <- y0 + rnorm(n^2,0.1)

# freq. lasso
las <- regress(diag(n^2),y-mean(y),'lasso')
lasCoef <- las$b[-1]





# horseshoe
hs <- bhs(X=diag(n^2),y,T=IT)

## t=100, m=46
## t=200, m=46
## t=300, m=54
```


```
## t=400, m=47
## t=500, m=50
## t=600, m=48
## t=700, m=51
## t=800, m=42
## t=900, m=56
```

```
hsPostMean <- colMeans(hs$beta[BI:IT,])
hsPostMed <- apply(hs$beta[BI:IT,],2,median)
par(mfrow=c(3,2))
par(pty="s")
image(matrix(y0,n),main='true',col=terrain.colors(100),zlim=c(-2,5))
image(matrix(y,n),main='observed',col=terrain.colors(100),zlim=c(-2,5))
image(matrix(lasCoef,n),main='freq lasso',col=terrain.colors(100),zlim=c(-2,5))
image(matrix(hsPostMean,n),main='horseshoe (mean)',col=terrain.colors(100),zlim=c(-2,5))
image(matrix(hsPostMed,n),main='horseshoe (med)',col=terrain.colors(100),zlim=c(-2,5))

#data.frame(y0,y,lasCoef,hsCoef)
```



**horseshoe (med)**



## ILLUSTRATION OF GIBBS SAMPLER FOR RIDGE REGRES-

## SION


First we call the bridge function, then we use some plots from the coda library. With this library trace plots of the mcmc sampler can easily be made.

```
IT <- 1000 # nr of mcmc iterations
BI <- 100  # nr of burnin iterations
```



```
d1 <- genSparseData(25,10,c(5, -3, 1, 0.2),1)
ridBayes <- bridge(d1$x,d1$y,T=IT,RJ=FALSE)
```

```
## t=100, m=10
## t=200, m=10
## t=300, m=10
## t=400, m=10
## t=500, m=10
## t=600, m=10
## t=700, m=10
## t=800, m=10
## t=900, m=10
```



```
plot(ridBayes, burnin=BI) # boxplots of coefs
```



```
plot(ridBayes, burnin=BI,which='s2') # plots for s2 (variance) chain
```









```
plot(ridBayes, burnin=BI, which='lambda2') # plots for lambda2 chain
```



```
library(coda)
```

```
mc.ridBayes.coef <- mcmc(ridBayes$beta[BI:IT,1:3], start=BI)
```


```
plot(mc.ridBayes.coef)
```



```
mc.ridBayes.s2 <- mcmc(ridBayes$s2[BI:IT], start=BI)
```


```
plot(mc.ridBayes.s2)
```

**Trace of var1**



Iterations


**Density of var1**



N = 901 Bandwidth = 0.08956


```
mc.ridBayes.lambda2 <- mcmc(ridBayes$lambda2[B1:IT], start=B1)  
plot(mc.ridBayes.lambda2)
```

**Trace of var1**



Iterations

**Density of var1**



N = 901 Bandwidth = 14.5